

Non-Traditional Online Training Methods for Always-On Models

Design goal: learning that emerges from information flow statistics and internal regulation rather than fixed dataset epochs or a single scalar objective.

1. Information-Theoretic Reinforcement

- adaptive adaptation where present latent improves future latent predictability.
- No external rewards. Objective is internal mutual information growth across small horizons.
- Use rolling-window estimators; keep compute O(1) per tick via EMA statistics.

Key equations:

Let z_t be latent state; estimate $I(z_t; z_{t+\Delta}) \approx H(z_{t+\Delta}) - H(z_{t+\Delta} | z_t)$.

Model H with Gaussian approximation: $H(z) \approx 0.5 \cdot \log \det(2\pi e \Sigma_z)$. Maintain Σ via exponential moving covariance.

Conditional entropy via linear predictor: $I_{\hat{A}}(z_t; z_{t+\Delta}) = A z_t$; residual $\epsilon = z_{t+\Delta} - I_{\hat{A}}$; $H(z_{t+\Delta} | z_t) \approx 0.5 \cdot \log \det(2\pi e \Sigma_\epsilon)$.

Signal $g_I = I_{\hat{A}}(z_t; z_{t+\Delta}) - \tau_I$ (target level τ_I). Plasticity gate: $\eta_{\text{eff}} = \text{clamp}(\eta_0 + k \cdot g_I, \eta_{\text{min}}, \eta_{\text{max}})$.

Online update sketch:

for each tick t : observe $z_t, z_{t+\Delta}$ update EMA covariance Σ_z and residual covariance Σ_ϵ from linear probe A (updated by least squares) $I_{\hat{A}} = 0.5 * (\log \det(\Sigma_z) - \log \det(\Sigma_\epsilon))$ $g_I = I_{\hat{A}} - \tau_I$ $\eta_{\text{eff}} = \text{clamp}(\eta_0 + k * g_I, \eta_{\text{min}}, \eta_{\text{max}})$ apply η_{eff} to parameter subsets that reduced Σ_ϵ this tick

2. Entropy■Regulated Flow Control

- Maintain a target entropy band for latent transitions to avoid collapse or chaos.
- Adjust plasticity, noise injection, and dropout to keep transition entropy in range.

Key equations:

$\Delta z_t = z_t - z_{t-1}$. Transition entropy $H_\Delta \approx 0.5 \cdot \log \det(2\pi e \Sigma_\Delta)$, with Σ_Δ EMA of Δz_t .

Homeostat $e = (H_\Delta - \mu_T) / \sigma_T$. Plasticity $\lambda = \sigma(\alpha \cdot e)$ mapped to $[\lambda_{\min}, \lambda_{\max}]$.

Noise scale $\sigma_{\text{noise}} = \sigma_0 \cdot \exp(-\beta \cdot e)$ (more noise when entropy is low).

Dropout $p_{\text{drop}} = \text{clamp}(p_0 + \gamma \cdot (e > 0 ? -e : |e|), 0, p_{\max})$.

Online update sketch:

per tick: $\Delta z = z_t - z_{t-1}$; update Σ_Δ $H_\Delta = 0.5 * \logdet(2\pi e \Sigma_\Delta)$ $e = (H_\Delta - \mu_T) / \sigma_T$ set learning-rate, noise, and dropout via $\lambda(e)$, $\sigma_{\text{noise}}(e)$, $p_{\text{drop}}(e)$

3. Predictive Symbiosis (Cross-Modal Co-Prediction)

Organisms co-predict each other, enforcing a shared manifold without labels.

Only update modules whose cross-prediction error improved over their own EMA.

Key equations:

Let $f_v(x_t) = z^v_t$, $g_a(y_t) = z^a_t$. Symmetric prediction: $z^{va}_{t+1} = P_{va}(z^a_t)$, $z^{av}_{t+1} = P_{av}(z^v_t)$.

Local tension $L_{sym} = ||z^v_{t+1} - z^{va}_{t+1}||^2 + ||z^a_{t+1} - z^{av}_{t+1}||^2$.

Gate by improvement $\delta = L_{av} - L_{sym}$. Update only if $\delta > 0$.

Online update sketch:

per tick with both modalities: compute z_v , z_a predict next states via P_{va} , P_{av} $L = mse(z_v_{next}, zhat_v) + mse(z_a_{next}, zhat_a)$ if $L < ema(L)$: allow weight updates for P_{va}/P_{av} and encoders; else freeze this tick

4. Self-Referential Distillation

- Spawn a shadow network S (small). Train S to imitate main network M online.
- Use divergence $D(M, S)$ as a structural growth signal; reconcile periodically.

Key equations:

Outputs o_M, o_S . Distill S by $L_{distill} = ||o_M - o_S||^2$.
Divergence signal $g_D = ||o_M - o_S||_2$ (or KL if probabilistic).
Plasticity gate for M: $\eta_{eff} = \eta_0 \cdot (1 + \kappa \cdot g_D)$ within bounds.

Online update sketch:

every K ticks: update $S \leftarrow \text{argmin } ||o_M - o_S||^2$ with small LR compute $g_D = \text{norm}(o_M - o_S)$ if $g_D > \tau$: increase plasticity or spawn a small adapter layer in M

5. Temporal Fractal Replay

- continuously replay stored compressed states at multiple time scales (1, 2, 4, 8... steps back).
- Adaptation occurs by aligning present dynamics with multiscale echoes, preventing drift.

Key equations:

Let C_k be a circular buffer with stride 2^k of compressed states c_t .

For each k , prediction residual $r_k = c_t - F(C_k[t])$. Minimize $\sum_k w_k ||r_k||^2$ but only if it reduces long-horizon variance.

Variance criterion: update if $\text{Var}_k^{\text{new}} < \text{Var}_k^{\text{EMA}} - \varepsilon$.

Online update sketch:

```
per tick: for k in {0..K}: fetch echo  $c_{t-2^k}$ ; predict forward via  $F$   $r_k = c_t - F(c_{t-2^k})$  if variance( $r_k$ ) improves EMA: do a tiny update on  $F$  and encoders
```

6. Internal Curiosity Economy

- Assign each latent channel i a plasticity budget b_i .
- Budget increases with novelty discovery and decreases with easy predictions.

Key equations:

Novelty $n_i = 1 - \cos(z_i, \mu_i)$ where μ_i is EMA centroid of channel i.

Budget update: $b_i \leftarrow \text{clamp}(b_i + \rho_n \cdot (n_i - \tau_n) - \rho_e \cdot \max(0, \tau_e - e_i), b_{\min}, b_{\max})$,
where e_i is per-channel prediction error, τ_n desired novelty, τ_e minimal error.

Effective LR per channel: $\eta_i = \eta_0 \cdot (b_i / b_{\max})$.

Online update sketch:

for channel i each tick: compute novelty n_i and error e_i update budget b_i scale
optimizer LR or gradient mask by η_i

7. Information-Equilibrium Decay (Synaptic Tagging)

bullet consolidates only when prediction error variance stabilizes, preventing overfitting to bursts.

bullet implements long-term retention without global checkpoints.

Key equations:

Track windowed mean μ_E and variance σ_E^2 of error E_t .

Stability score $s = 1 / (1 + \sigma_E^2)$. Consolidate when $s > s_{thr}$ and $|\mu_E - \mu_E^{(EMA)}| < \delta$.

Consolidation operator: $\theta \leftarrow (1 - \lambda)\theta + \lambda \cdot \theta'$ where θ' are weights after micro-updates; λ depends on s .

Online update sketch:

per window W : compute μ_E , σ_E^2 if stability high: blend current weights into a slow-moving consolidated copy otherwise: keep learning fast but do not consolidate

8. Signal Resonance Training

- bullet Detect resonant structure between input dynamics and internal activations in frequency space.
- bullet Strengthen connections that co-oscillate; weaken those that remain phase-incoherent.

Key equations:

Compute short FFTs: $X_f = \text{FFT}(x_{\{t-L:t\}})$, $Z_f = \text{FFT}(z_{\{t-L:t\}})$.

Cross-power spectral density: $S_{xz}(f) = X_f \cdot \text{conj}(Z_f)$. Coherence $C(f) = |S_{xz}(f)|^2 / (S_{xx}(f) S_{zz}(f))$.

Resonance score $R = \sum_f w_f \cdot C(f)$. Update gate $g_R = R - \tau_R$.

Weight change $\Delta W \propto g_R \cdot (\partial z / \partial W)$ projected onto frequencies with highest coherence.

Online update sketch:

every L ticks: compute FFT windows for input and latent measure coherence $C(f)$; get resonance R modulate gradients by g_R concentrating updates on resonant pathways

Integration Sketch

Combine methods via gates. At each tick, compute signals {g_I, e, δ , g_D, Var_k, b_i, s, g_R}. Map them to per-module plasticity multipliers and gradient masks. Ensure O(1) memory per method using EMA statistics and bounded buffers.

Notation: logdet denotes the logarithm of the matrix determinant; Σ terms are covariance estimates tracked online via exponential moving averages; clamp() limits to a range.